

ASSIGNMENT-03

Build CNN Model for Classification Of Flowers

Assignment Date	5 October 2022
Student Name	K.S.Naveen kumar
Student Roll Number	113219071025
Maximum Marks	2 Marks

QUESTION 1:

Download the Dataset

Dataset is downloaded and uploaded

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import tensorflow as tf

from keras.models import Sequential
from keras.layers import Convolution2D
from keras.layers import MaxPooling2D
from keras.layers import Flatten
from keras.layers import Dense

from tensorflow.keras.preprocessing import image
from tensorflow.keras.preprocessing.image import ImageDataGenerator

import cv2
```

QUESTION 2:

Image Augmentation

```
data_path = '/content/drive/MyDrive/CNN/flowers'
```

```
batch size = 32
```

```
target_size = (64, 64)
```

```
train_datagen = ImageDataGenerator(rescale=1./255,
                                   shear_range=0.2,
                                   zoom_range=0.2,
                                   width_shift_range=0.1,
                                   height_shift_range=0.1,
                                   horizontal_flip=True,
                                   validation_split=0.2)
```

```
test datagen = ImageDataGenerator(rescale=1. / 255, validation split=0.2)
```

[illegible]

```
subset="training",
class_mode = 'categorical')
```

```
X_test = test_datagen.flow_from_directory(data_path,
                                         target_size=target_size,
                                         batch_size=batch_size,
                                         subset="validation",
                                         class_mode='categorical')
```

```
[2] data_path = '/content/drive/MyDrive/CNN/flowers'
     batch_size = 32
     target_size = (64, 64)

[3] train_datagen = ImageDataGenerator(rescale=1./255,
                                       shear_range=0.2,
                                       zoom_range=0.2,
                                       width_shift_range=0.1,
                                       height_shift_range=0.1,
                                       horizontal_flip=True,
                                       validation_split=0.2)

test_datagen = ImageDataGenerator(rescale=1. / 255, validation_split=0.2)

[4] X_train = train_datagen.flow_from_directory(data_path,
                                              target_size=target_size,
                                              batch_size=batch_size,
                                              subset="training",
                                              class_mode = 'categorical')

X_test = test_datagen.flow_from_directory(data_path,
                                          target_size=target_size,
                                          batch_size=batch_size,
                                          subset="validation",
                                          class_mode='categorical')

Found 16 images belonging to 5 classes.
Found 4 images belonging to 5 classes.
```

QUESTION 3:

Create Model

```
model = Sequential()
```

```
+ Code + Text

[5] model = Sequential()
```

QUESTION 4:

Add Layers (Convolution,MaxPooling,Flatten,Dense-(Hidden Layers),Output

```
model.add(Convolution2D(32, (3, 3), input_shape=(64, 64, 3),
activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

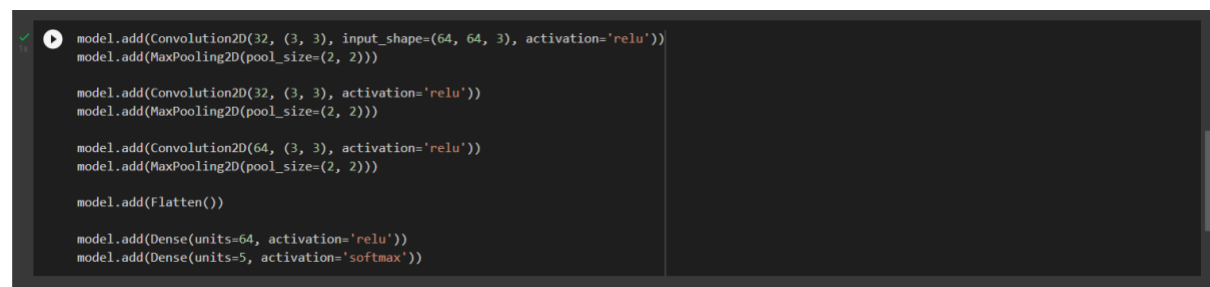
model.add(Convolution2D(32, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Convolution2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())

model.add(Dense(units=64, activation='relu'))
model.add(Dense(units=5, activation='softmax'))

model.summary()
```

A screenshot of a Jupyter Notebook cell with a green checkmark icon on the left. The cell contains the same Python code as shown in the previous block, defining a sequential model with convolutional, pooling, flattening, and dense layers.

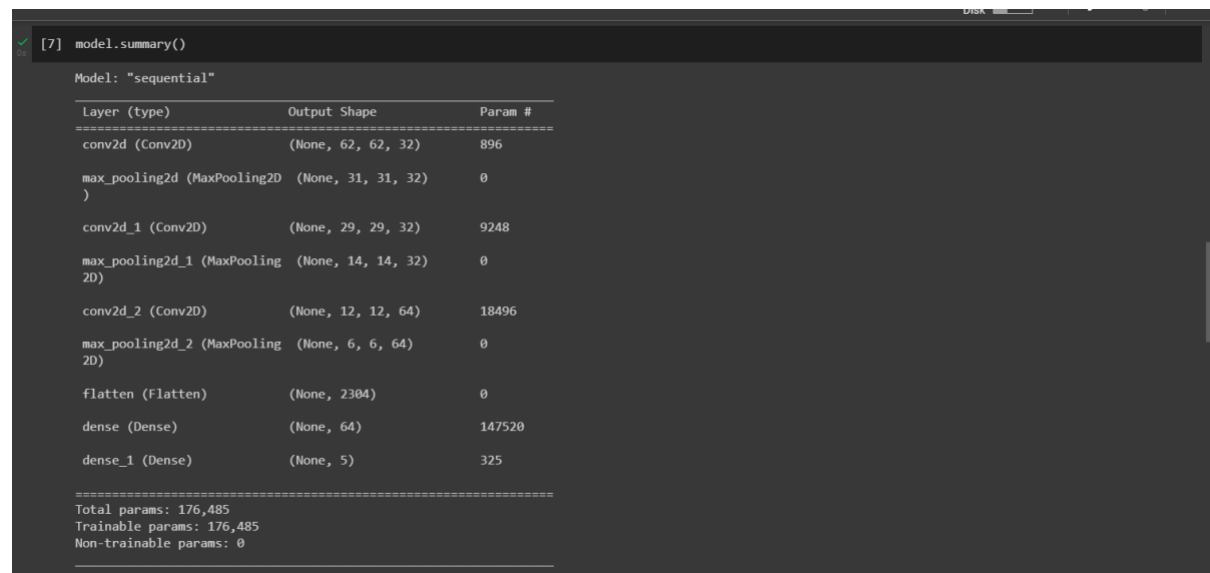
```
model.add(Convolution2D(32, (3, 3), input_shape=(64, 64, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Convolution2D(32, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Convolution2D(64, (3, 3), activation='relu'))
model.add(MaxPooling2D(pool_size=(2, 2)))

model.add(Flatten())

model.add(Dense(units=64, activation='relu'))
model.add(Dense(units=5, activation='softmax'))
```

A screenshot of a Jupyter Notebook cell showing the output of the model.summary() command. The output is a table with three columns: Layer (type), Output Shape, and Param #. Below the table, it shows the total number of parameters, trainable parameters, and non-trainable parameters.

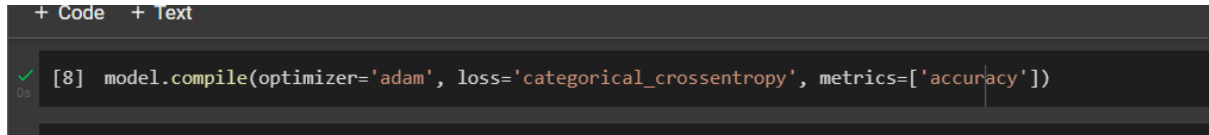
```
[7] model.summary()

Model: "sequential"
=====
Layer (type)                Output Shape                Param #
-----
conv2d (Conv2D)              (None, 62, 62, 32)         896
max_pooling2d (MaxPooling2D) (None, 31, 31, 32)         0
conv2d_1 (Conv2D)            (None, 29, 29, 32)         9248
max_pooling2d_1 (MaxPooling2D) (None, 14, 14, 32)         0
conv2d_2 (Conv2D)            (None, 12, 12, 64)         18496
max_pooling2d_2 (MaxPooling2D) (None, 6, 6, 64)          0
flatten (Flatten)            (None, 2304)               0
dense (Dense)                (None, 64)                 147520
dense_1 (Dense)              (None, 5)                  325
=====
Total params: 176,485
Trainable params: 176,485
Non-trainable params: 0
```

QUESTION 5:

Compile the Model

```
model.compile(optimizer='adam', loss='categorical_crossentropy',  
metrics=['accuracy'])
```

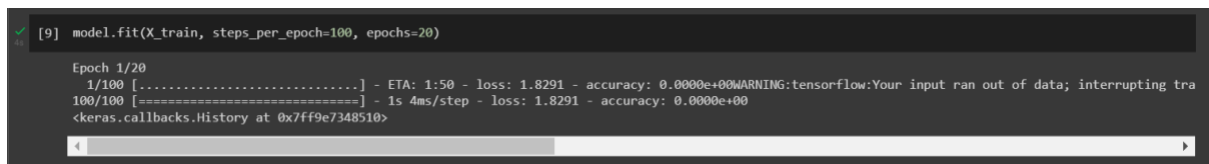


A Jupyter Notebook cell with a dark background. At the top, there are tabs for '+ Code' and '+ Text'. Below the tabs, a green checkmark icon and '0s' are visible. The code cell contains the following line: `[8] model.compile(optimizer='adam', loss='categorical_crossentropy', metrics=['accuracy'])`.

QUESTION 6:

Fit the model

```
model.fit(X_train, steps_per_epoch=100, epochs=20)
```

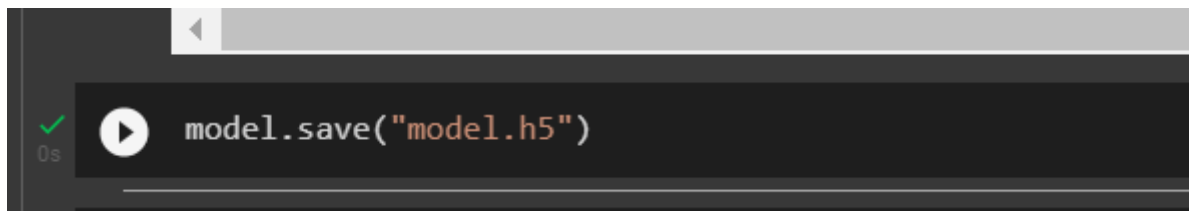


A Jupyter Notebook cell with a dark background. At the top, a green checkmark icon and '0s' are visible. The code cell contains the following line: `[9] model.fit(X_train, steps_per_epoch=100, epochs=20)`. Below the code, the output shows the progress of the training:
Epoch 1/20
1/100 [.....] - ETA: 1:50 - loss: 1.8291 - accuracy: 0.0000e+00
100/100 [=====] - 1s 4ms/step - loss: 1.8291 - accuracy: 0.0000e+00
<keras.callbacks.History at 0x7ff9e7348510>

QUESTION 7:

Save the Model

```
model.save("model.h5")
```



A Jupyter Notebook cell with a dark background. At the top, a green checkmark icon and '0s' are visible. The code cell contains the following line: `model.save("model.h5")`. Below the code, there is a play button icon.

QUESTION 8:

Test the Model

```
def predict():  
    img =  
    image.load_img("/content/drive/MyDrive/CNN/flowers/rose/118974357_0faa23cce  
9_n.jpg", target_size=target_size)  
    x = image.img_to_array(img)  
    x = tf.expand_dims(x,0)  
  
    labels = ['daisy', 'dandelion', 'rose', 'sunflower', 'tulip']  
  
    pred = model.predict(x)  
    prediction = labels[np.argmax(pred[0])]
```

```

print(f'The given image is a {prediction}')

plt.imshow(plt.imread("/content/drive/MyDrive/CNN/flowers/rose/118974357_0f
aa23cce9_n.jpg"))
plt.axis('off')
plt.show()

predict()

```

```

[11] def predict():
    img = image.load_img("/content/drive/MyDrive/CNN/flowers/rose/118974357_0faa23cce9_n.jpg", target_size=target_size)
    x = image.img_to_array(img)
    x = tf.expand_dims(x,0)

    labels = ['daisy', 'dandelion', 'rose', 'sunflower', 'tulip']

    pred = model.predict(x)
    prediction = labels[np.argmax(pred[0])]

    print(f'The given image is a {prediction}')
    plt.imshow(plt.imread("/content/drive/MyDrive/CNN/flowers/rose/118974357_0faa23cce9_n.jpg"))
    plt.axis('off')
    plt.show()

```

